

TAIXING BI

Senior AI Infrastructure & Platform Engineer

LLM Systems • AI Platform • Kubernetes • GPU Inference

New York, NY (Open to Remote / Hybrid / NYC)

347-560-4308 • taixingbijob@gmail.com

LinkedIn: [linkedin.com/in/taixing-bi-757a4961](https://www.linkedin.com/in/taixing-bi-757a4961)

Portfolio: <https://taixingai.com>

GitHub: github.com/taixingbi

SUMMARY

Senior AI Infrastructure & Platform Engineer with 7+ years of experience designing, building, and operating production AI systems, distributed platforms, and cloud-native infrastructure.

Specialized in GPU inference, LLM serving, Retrieval-Augmented Generation (RAG), agent orchestration, evaluation frameworks, and Kubernetes-based AI platforms. Proven track record delivering end-to-end AI systems spanning model serving, retrieval, routing, evaluation, observability, and production operations.

Experienced across AI infrastructure, distributed systems, cloud-native platforms, backend engineering, and production reliability. Strong background building scalable services supporting machine learning, analytics, operational intelligence, and enterprise AI workloads.

CORE SKILLS

AI Platform & Infrastructure

Kubernetes • Docker • ArgoCD • GitOps • CI/CD • Distributed Systems • Platform Engineering • Service Reliability

AI Serving & Inference

vLLM • SGLang • Hugging Face TGI • LoRA Serving • Model Routing • Token Streaming • Batching • Caching • GPU Workload Management

AI Systems

Retrieval-Augmented Generation (RAG) • Agent Orchestration • Tool Calling • LangChain • LangGraph • Evaluation Frameworks • Fine-Tuning (LoRA, SFT, DPO)

Cloud & Data Infrastructure

AWS (Lambda, DynamoDB, S3, EC2, SQS, EventBridge) • Kafka / MSK • Snowflake

Observability

Prometheus • Grafana • Datadog • CloudWatch • Distributed Tracing • Performance Monitoring

Programming

Python • TypeScript • Java • SQL

SELECTED AI PLATFORM EXPERIENCE

HuntAI — AI Platform & Infrastructure

Founder & Architect | <https://taixingai.com>

- Designed, built, and operated an end-to-end AI platform supporting LLM serving, retrieval, orchestration, evaluation, observability, and GitOps deployment workflows.
- Architected and implemented a Kubernetes-based AI platform consisting of 10+ production services spanning gateway, orchestration, retrieval, model serving, evaluation, monitoring, and deployment automation.
- Built AI Gateway services supporting model routing, tool execution, authentication, streaming responses, usage tracking, and observability.
- Designed and operated multi-GPU inference infrastructure using vLLM, Qwen, and LoRA adapters for production AI serving workloads.
- Implemented hybrid retrieval pipelines combining semantic search, keyword search, reranking, citation grounding, and retrieval evaluation.
- Developed orchestration services coordinating routing, retrieval, tool execution, validation, and response generation workflows.
- Built automated evaluation frameworks using golden datasets, routing accuracy testing, regression validation, and model quality monitoring.

- Designed GitOps deployment workflows using ArgoCD and GitHub Actions supporting automated deployments, environment promotion, and drift detection.
- Built observability platforms using Prometheus and Grafana providing centralized monitoring, metrics, dashboards, and operational visibility.
- Fine-tuned and evaluated routing models using LoRA, SFT, and DPO workflows with automated benchmarking and regression testing pipelines.

Selected Platform Metrics

10+ production microservices running on Kubernetes

~100ms chat time-to-first-token (TTFT)

180+ RPS embedding service throughput

200+ RPS reranking service throughput

6–7 sustained generation requests per second

Technologies: Kubernetes, ArgoCD, vLLM, Qwen, Python, TypeScript, FastAPI, Qdrant, Prometheus, Grafana, Cloudflare, GitHub Actions

PROFESSIONAL EXPERIENCE

Saks.com

Senior Software Engineer / AI Platform Engineer

New York, NY | Aug 2021 – Present

- Led the design and implementation of production AI platform capabilities supporting retrieval, operational intelligence, engineering productivity, and cloud-native services.
- Developed Retrieval-Augmented Generation (RAG) systems including document ingestion, chunking, embedding generation, hybrid retrieval, reranking, and citation grounding.
- Built AI-powered operational intelligence applications enabling natural-language access to operational, engineering, and business knowledge sources through retrieval and tool-calling workflows.
- Architected agent orchestration workflows coordinating retrieval, reasoning, validation, and response generation using graph-based execution models.
- Designed AI evaluation frameworks supporting regression testing, prompt versioning, model comparison, automated validation, and production quality monitoring.
- Implemented low-latency inference and streaming response services with distributed tracing, observability, monitoring, and operational controls.
- Built cloud-native AI services leveraging Kubernetes, AWS, Kafka, DynamoDB, Lambda, and event-driven architectures.
- Developed centralized data and analytics platforms supporting AI, machine learning, operational reporting, and retrieval workloads.
- Implemented production reliability, observability, and monitoring solutions using Datadog, CloudWatch, Grafana, and automated alerting systems.
- Partnered with engineering leadership and cross-functional teams to identify, design, and deliver AI platform capabilities from concept through production deployment.

Technologies: AWS, Kubernetes, Kafka, DynamoDB, Lambda, Snowflake, Python, Java, TypeScript, Datadog, CloudWatch

Kaden Health

Machine Learning Engineer

Remote | Jul 2019 – Aug 2021

- Developed NLP and machine learning solutions for healthcare applications using TensorFlow, PyTorch, and spaCy.
- Built end-to-end machine learning pipelines covering data processing, feature engineering, model training, evaluation, deployment, and monitoring.
- Designed and implemented production inference services and model-serving APIs supporting real-time clinical workflows.
- Developed text classification, sentiment analysis, entity extraction, and predictive modeling solutions for healthcare datasets.
- Built automated evaluation and validation pipelines improving model quality, reliability, and deployment readiness.

- Designed event-driven data processing architectures using Kafka and NATS supporting real-time machine learning workloads.
 - Built centralized analytics, logging, and monitoring pipelines using BigQuery and cloud-native infrastructure.
 - Collaborated with product, engineering, and clinical stakeholders to deliver production AI capabilities.
- Technologies: Python, TensorFlow, PyTorch, spaCy, Kafka, NATS, BigQuery, GCP

CGMAX

Machine Learning Engineer

New York, NY | Jan 2018 – Jul 2019

- Developed machine learning solutions for financial forecasting, sentiment analysis, and quantitative research applications.
- Built predictive modeling pipelines using structured and unstructured financial datasets to generate investment and market insights.
- Designed distributed feature engineering and data processing platforms using Spark, Hadoop, Kafka, and HBase.
- Developed real-time data ingestion and analytics systems supporting machine learning and quantitative research workloads.
- Built model evaluation, validation, and backtesting frameworks improving prediction quality and deployment readiness.
- Partnered with quantitative researchers and software engineers to operationalize machine learning capabilities in production environments.

Technologies: Python, Spark, Hadoop, Kafka, HBase, SQL

NVIDIA

Software Engineer

Santa Clara, CA | Mar 2016 – Dec 2017

- Developed telemetry and observability platforms supporting GPU engineering teams.
- Built backend services and APIs for collecting and analyzing GPU performance, utilization, and diagnostic data.
- Designed analytics dashboards enabling engineers to investigate system behavior and performance trends at scale.
- Developed engineering tools and internal applications using React, Python, and Node.js.
- Partnered with cross-functional hardware and software teams to improve visibility into GPU performance and operational health.
- Gained early experience working with large-scale GPU systems, performance analytics, observability, and infrastructure monitoring.

Technologies: Python, Node.js, React, Analytics Platforms, Monitoring Systems

EDUCATION

Master of Science, Analytics (Machine Learning & NLP)

Harrisburg University

2019 – 2022

Master of Science, Electrical & Computer Engineering

University of Alabama in Huntsville

2013 – 2015

PUBLICATION

A Novel Approach to Impedance-Based Fault Location for High Voltage Cables

IEEE Industry Applications Society (IAS)